

The Stata Journal (2015)
15, Number 2, pp. 480–500

Conducting interrupted time-series analysis for single- and multiple-group comparisons

Ariel Linden
Linden Consulting Group, LLC
Ann Arbor, MI
alinden@lindenconsulting.org

Abstract. In this article, I introduce the `itsa` command, which performs interrupted time-series analysis for single- and multiple-group comparisons. In an interrupted time-series analysis, an outcome variable is observed over multiple, equally spaced time periods before and after the introduction of an intervention that is expected to interrupt its level or trend. The `itsa` command estimates the effect of an intervention on an outcome variable either for a single treatment group or when compared with one or more control groups. Additionally, its options allow the user to control for autocorrelated disturbances and to estimate treatment effects over multiple periods.

Keywords: st0389, `itsa`, interrupted time series, quasi-experimental designs, causal inference

1 Introduction

In considering the impact of large-scale interventions (for example, population-based health interventions, media campaigns, and dissemination of professional guidelines) or public policy changes (for example, new laws or taxes), researchers are often faced with an effective sample size of $N = 1$, where the treated group may be the local community, state, or an even larger unit. It is also fairly common in these situations that the only data available are reported at an aggregate level (for example, morbidity or mortality rates, average costs, and median incomes). If multiple observations on an outcome variable of interest in the preintervention and postintervention periods can be obtained, an interrupted time-series analysis (ITSA) offers a quasi-experimental research design with a potentially high degree of internal validity (Campbell and Stanley 1966; Shadish, Cook, and Campbell 2002). Naturally, when the treated group's outcomes can also be contrasted with those of one or more comparison groups, the internal validity is further enhanced by allowing the researcher to potentially control for confounding omitted variables.

ITSA has been used in many areas of study, such as assessing the effects of community interventions (Biglan, Ary, and Wagenaar 2000; Gillings, Makuc, and Siegel 1981), public policy (Muller 2004), regulatory actions (Briesacher et al. 2013), and health technology assessment (Ramsay et al. 2003), to name but a few. ITSA has also been proposed as a more flexible and rapid design to be considered in health research before defaulting to the traditional two-arm randomized controlled trial (Riley et al. 2013). In addition,

systematic reviews of the literature are increasingly including studies that have used ITSA as their primary research design (Cochrane Effective Practice and Organisation of Care [EPOC] 2013).

In this article, I introduce the new `itsa` command, which performs interrupted time-series analysis using two ordinary least-squares (OLS) regression-based approaches available in the official Stata packages `newey` and `prais`. Additionally, `itsa` can estimate treatment effects for multiple treatment periods.

2 Method and formulas

Statistical analyses used for ITSA must account for autocorrelated data. The two general approaches historically used in ITSA are autoregressive integrated moving-average models (see Box and Tiao [1975], Glass, Willson, and Gottman [1975], and McDowall et al. [1980]) and OLS regression models designed to adjust for autocorrelation (see, among others, Crosbie [1993]; Gottman [1981]; McKnight, McKean, and Huitema [2000]; Simonton [1977a]; and Velicer and McDonald [1991]). `itsa` relies on OLS rather than on regression methods based on autoregressive integrated moving-average models because the former is often more flexible and broadly applicable in an interrupted time-series context (Box and Jenkins 1976; Velicer and Harrop 1983).

2.1 The single-group analysis

When there is only one group under study (no comparison groups), the standard ITSA regression model assumes the following form (Huitema and McKean 2000a; Linden and Adams 2011; Simonton 1977a; Simonton 1977b):

$$Y_t = \beta_0 + \beta_1 T_t + \beta_2 X_t + \beta_3 X_t T_t + \epsilon_t \quad (1)$$

Y_t is the aggregated outcome variable measured at each equally spaced time point t , T_t is the time since the start of the study, X_t is a dummy (indicator) variable representing the intervention (preintervention periods 0, otherwise 1), and $X_t T_t$ is an interaction term. These terms are displayed in the lower half of figure 1. In the case of a single-group study, β_0 represents the intercept or starting level of the outcome variable. β_1 is the slope or trajectory of the outcome variable until the introduction of the intervention. β_2 represents the change in the level of the outcome that occurs in the period immediately following the introduction of the intervention (compared with the counterfactual). β_3 represents the difference between preintervention and postintervention slopes of the outcome. Thus we look for significant p -values in β_2 to indicate an immediate treatment effect, or in β_3 to indicate a treatment effect over time (Linden and Adams 2011).




Figure 1. Visual depiction of a single group (lower line) and multiple group (upper and lower lines) interrupted time-series design, from [Linden and Adams \(2011\)](#)

Legend:

Single group— β_0 : intercept; β_1 : slope prior to intervention; β_2 : change in level in the period immediately following intervention initiation (compared with counterfactual); β_3 : difference between preintervention and postintervention slopes.

Multiple group— β_0 to β_3 represent the control group; β_4 to β_7 represent the treatment group. β_4 : difference in the level between treatment and control prior to intervention; β_5 : difference in the slope between treatment and control prior to intervention; β_6 : difference in the level between treatment and control in the period immediately following intervention initiation; β_7 : difference between treatment and control in the slope after initiation of the intervention compared with preintervention.

When the random error terms follow a first-order autoregressive [AR(1)] process,

$$\epsilon_t = \rho\epsilon_{t-1} + u_t \quad (2)$$

where the autocorrelation parameter ρ is the correlation coefficient between adjacent error terms, such that $|\rho| < 1$, and the disturbances u_t are independent $N(0, \sigma^2)$ (see Kutner et al. [2005] for a detailed discussion of autocorrelation in time-series models).

Identification in both the single- and multiple-group models is driven by the functional-form assumptions of the ITSA model. By design, a single-group ITSA has no comparable control group; rather, the preintervention trend projected into the treatment period serves as the counterfactual. We assume that any time-varying unmeasured

confounder is relatively slowly changing so that it would be distinguishable from the sharp jump of the intervention indicator. This underscores the need for caution with these methods if there are multiple policy shifts in the time window around the implementation of the intervention.

The assumptions necessary for causal inference in the single-group ITSA may seem plausible when the preintervention trend is flat followed by a significant change in the outcome variable immediately following the introduction of the intervention and then sustained over time. However, these assumptions may seem less plausible if a trend already exists in the time series prior to the intervention. While the ITSA literature does not address the topic of testing for interruptions in the level and trend of the outcome variable [β_2 and β_3 of (1)] prior to the actual period in which the intervention started, we can look to the regression-discontinuity literature to provide guidance for applicable robustness tests. In practice, this would simply entail testing for interruptions after replacing the true intervention start period with other pseudo-start periods along the preintervention continuum.

In an adaptation of [Imbens and Lemieux \(2008\)](#) for ITSA, an investigator could use the median time point of the preintervention period to test for an interruption. In a sufficiently long time series, the median time point of the preintervention period is a good choice of a pseudo-start period to maximize power to detect a significant jump (because the subsample will be evenly split on both sides). For shorter time series, a simple iterative process of testing each preintervention time period as the pseudo-start period may be a good approach. In using such robustness tests, the underlying assumptions of the single-group ITSA may be challenged if interruptions in the level or trend of the outcome variable are found to exist at other time points prior to the true initiation of the intervention.

2.2 The multiple-group analysis

When one or more control groups are available for comparison, the regression model in (1) is expanded to include four additional terms (β_4 to β_7) ([Linden and Adams 2011](#); [Simonton 1977a](#); [Simonton 1977b](#)):

$$Y_t = \beta_0 + \beta_1 T_t + \beta_2 X_t + \beta_3 X_t T_t + \beta_4 Z + \beta_5 Z T_t + \beta_6 Z X_t + \beta_7 Z X_t T_t + \epsilon_t \quad (3)$$

Here Z is a dummy variable to denote the cohort assignment (treatment or control), and $Z T_t$, $Z X_t$, and $Z X_t T_t$ are all interaction terms among previously described variables. Now when examining figure 1, the coefficients of the lower line, β_0 to β_3 , represent the control group, and the coefficients of the upper line, β_4 to β_7 , represent values of the treatment group. More specifically, β_4 represents the difference in the level (intercept) of the outcome variable between treatment and controls prior to the intervention, β_5 represents the difference in the slope (trend) of the outcome variable between treatment and controls prior to the intervention, β_6 indicates the difference between treatment and control groups in the level of the outcome variable immediately following introduction of the intervention, and β_7 represents the difference between treatment and control groups in the slope (trend) of the outcome variable after initiation of the intervention compared

with preintervention (akin to a difference-in-differences of slopes). If the multiple-group model follows an AR(1) process, the random error term is defined as in (2).

A multiple-group ITSA may be particularly valuable when there is an exogenous policy shift that affects all the groups. The key assumption is that the change in the level or trend in the outcome variable is presumed to be the same both for the control group and, counterfactually, for the treatment group had it not received the intervention. In other words, we assume that confounding omitted variables affect both treatment and control groups similarly. A major strength of the multiple-group ITSA is the ability to test for comparability between groups on observed covariates and in particular, the two parameters β_4 and β_5 , which play a particularly important role in establishing whether the treatment and control groups are balanced on both the level and the trajectory of the outcome variable in the preintervention period. If these data were from a randomized controlled trial, we would expect similar levels and slopes prior to the intervention. However, in an observational study where equivalence between groups cannot be ensured, any observed differences will likely raise concerns about the ability to draw causal inferences about the relationship between the intervention and the outcomes (Linden and Adams 2011).

To reduce the threat of confounding, investigators may attempt to emulate the randomization process with observational data by finding control groups that are comparable to the treatment group on observed preintervention covariates. One approach to finding comparable controls out of a pool of potential candidates is via an iterative process in which each non-treated group is compared separately with the treatment group using the model defined in (3). Those groups who have p -values greater than 0.05 (or a higher threshold) on both β_4 and β_5 can be selected as controls for inclusion in the final model. This method can be easily expanded to other available covariates; however, there is a diminished likelihood of finding good controls as the number of covariates is increased. If achieving balance on many covariates is an important factor, two alternative approaches to *itsa* should be considered: the synthetic controls approach described by Abadie, Diamond, and Hainmueller (2010) and implemented in Stata using the *synth* package (Abadie, Diamond, and Hainmueller 2014), or the propensity-score weighting technique described by Linden and Adams (2011).

2.3 Data variables corresponding to model parameters

Table 1 displays the variables used in regression models (1) and (3), using an artificial example with one intervention period. There are two individuals (or groups) in these data ($ID = 1, 2$) with six observations each (T). X indicates that there are two preintervention observations, followed by four observations in the intervention period (the intervention commences when $T = 3$). XT is an interaction term of $X \times T$, which starts in the observation period immediately following the start of the intervention ($T = 4$) and runs sequentially until the last observation when $T = 6$ (see Huitema and McKean [2000a] for an exposition on the appropriateness of commencing the sequence in the observation period after the start of the intervention). Here we transform $XT = (T - 3) \times X$ so that it runs sequentially starting at 1. Additional variables are required for a multiple-group

analysis. Z indicates the treatment status, where $Z = 1$ for the treatment group and $Z = 0$ for the control group. ZT , ZX , and ZXT are additional interaction terms used in multiple-group comparisons, as described above in section 2.2. When multiple treatment periods are specified, additional variables are added to the dataset, corresponding to each treatment period respectively (see section 4.3 for an example).

Upon running the `itsa` command, all variables required for the corresponding single- or multiple-group model are automatically generated and added to the dataset.

Table 1. Covariates used in a single-group ITSA (T , X , XT) and multiple-group ITSA (T , X , XT , Z , ZT , ZX , ZXT) corresponding to regression models (1) and (3), respectively

ID	T	X	XT	Z	ZT	ZX	ZXT
1	1	0	0	1	1	0	0
1	2	0	0	1	2	0	0
1	3	1	0	1	3	1	0
1	4	1	1	1	4	1	1
1	5	1	2	1	5	1	2
1	6	1	3	1	6	1	3
2	1	0	0	0	0	0	0
2	2	0	0	0	0	0	0
2	3	1	0	0	0	0	0
2	4	1	1	0	0	0	0
2	5	1	2	0	0	0	0
2	6	1	3	0	0	0	0

2.4 Models

`itsa` allows the user to choose between two OLS regression-based models specifically designed for time-series data. The first, **newey** (see [TS] **newey**), estimates the coefficients by OLS regression but produces Newey–West standard errors to handle autocorrelation in addition to possible heteroskedasticity. The second model, **prais** (see [TS] **prais**), uses the generalized least-squares method to estimate the parameters in a linear regression model in which the errors are assumed to follow an AR(1) process. More specifically, **prais** offers several methods to transform the original observations based on the pooled autocorrelation estimate ρ to remove the correlation between first-order errors (that is, the correlation between the errors of each observation period and those of the preceding observation period).

The type of model that an investigator will choose for conducting time-series analysis will likely depend on a combination of factors, with primary attention on the number of lags in the data for which autocorrelation is present. In general, the investigator first fits an OLS model using either **regress** or **newey** (with **lag(0)** specified) and then tests for autocorrelation in the error distribution. It is important to test for the presence of autocorrelated errors when using regression-based time-series methods, because such tests provide critical diagnostic information regarding the adequacy of the time-series model (that is, whether tests and confidence intervals on the regression coefficients are satisfactory, whether important variables have been left out of the time-series regression model and because autocorrelated errors are produced when the functional form of the variables included in the model is incorrect; Huitema and McKean [2000b]). The package offers several postestimation commands for this purpose (see [R] **regress postestimation time series**). In addition, there is a comprehensive and versatile user-written program, **actest** (Baum and Schaffer 2013), that is downloadable from the Statistical Software Components archive, with the default being the Cumby–Huizinga general test for autocorrelation (Cumby and Huizinga 1992).

3 The **itsa** command

3.1 Syntax

```
itsa depvar [indepvars] [if] [in] [weight], trperiod(numlist) [single
treatid(#) contid(numlist) prais lag(#) figure posttrend replace
prefix(string) model_options]
```

A dataset for a single panel must be declared to be time-series data using **tsset** *timevar*. When the dataset contains multiple panels, a strongly balanced panel dataset using **tsset** *panelvar timevar* must be declared; see [TS] **tsset**. *indepvars* may contain factor variables; see [U] 11.4.3 **Factor variables**. *depvar* and *indepvars* may contain time-series operators; see [U] 11.4.4 **Time-series varlists**. **aweights** are allowed; see [U] 11.1.6 **weight**. See [TS] **newey postestimation** and [TS] **prais postestimation** for features available after estimation.

3.2 Options

trperiod(*numlist*) specifies the time period when the intervention begins. The values entered for the time period must be in the same units as the panel time variable specified in **tsset** *timevar*; see [TS] **tsset**. More than one period may be specified. **trperiod()** is required.

single indicates that **itsa** will be used for a single-group analysis. Conversely, omitting **single** indicates that **itsa** is for a multiple-group comparison.

treatid(#) specifies the identifier of the single treated unit under study when the dataset contains multiple panels. The value entered must be in the same units as the panel variable specified in **tsset** *panelvar timevar*; see [TS] **tsset**. When the dataset contains data for only a single panel, **treatid()** must be omitted.

contid(numlist) specifies a list of identifiers to be used as control units in the multiple-group analysis. The values entered must be in the same units as the panel variable specified in **tsset** *panelvar timevar*; see [TS] **tsset**. If **contid()** is not specified, all nontreated units in the data will be used as controls.

prais specifies to fit a **prais** model. If **prais** is not specified, **itsa** will use **newey** as the default model.

lag(#) specifies the maximum lag to be considered in the autocorrelation structure when a **newey** model is chosen. If the user specifies **lag(0)**, the default, the output is the same as **regress, vce(robust)**. An error message will appear if both **prais** and **lag()** are specified, because **prais** implements an AR(1) model by design.

figure produces a line plot of the predicted *depvar* variable combined with a scatterplot of the actual values of *depvar* over time. In a multiple-group analysis, **figure** plots the average values of all controls used in the analysis (more specifically, data for specified controls are collapsed and the monthly observations are averaged).

posttrend produces posttreatment trend estimates using **lincom**, for the specified model. In the case of a single-group ITSA, one estimate is produced. In the case of a multiple-group ITSA, an estimate is produced for the treatment group, the control group, and the difference. In the case of multiple treatment periods, a separate table is produced for each treatment period.

replace replaces variables created by **itsa** if they already exist. If **prefix()** is specified, only variables created by **itsa** with the same prefix will be replaced.

prefix(string) adds a prefix to the names of variables created by **itsa**. Short prefixes are recommended.

model_options specify all available options for **prais** when the **prais** option is chosen; otherwise, all available options for **newey** other than **lag()** are specified.

3.3 Stored results

Because **itsa** passes all user-entered information to **prais** and **newey**, all results stored by those commands are available. Additionally, **itsa** generates several key time-series variables and adds them to the current dataset, as described in section 2.3. These additional variables allow the user to further estimate treatment effects using **arma** or other time-series models.

Table 2 is a cross reference to default names for those variables that appear in the regression output tables (and used when **posttrend** is specified). Variables starting with **_z** are added to the dataset only when a multiple-group comparison is specified.

(`trperiod`) is a suffix added to certain variables indicating the start of the intervention period. This is particularly helpful for differentiating between added variables when multiple interventions are specified (see the example presented in section 4.3). If the user specifies a `prefix()`, it will be applied to all variables generated by `itsa`.

Table 2. Descriptions of default names for variables that appear in the regression output tables

Variable	Description
<code>_t</code>	time since start of study
<code>_x(trperiod)</code>	dummy variable representing the intervention periods (preintervention periods 0, otherwise 1)
<code>_x.t(trperiod)</code>	interaction of <code>_x</code> and <code>_t</code>
<code>_z</code>	dummy variable to denote the cohort assignment (treatment or control)
<code>_z.x(trperiod)</code>	interaction of <code>_z</code> and <code>_x</code>
<code>_z.x.t(trperiod)</code>	interaction of <code>_z</code> , <code>_x</code> , and <code>_t</code>
<code>_s_depvar_pred</code>	predicted value generated after running <code>itsa</code> for a single group
<code>_m_depvar_pred</code>	predicted value generated after running <code>itsa</code> for a multiple-group comparison

4 Examples

In 1988, California passed the voter-initiative Proposition 99, which was a widespread effort to reduce smoking rates by raising the cigarette excise tax by 25 cents per pack and to fund anti-smoking campaigns and other related activities throughout the state (for a comprehensive discussion of this initiative, see [Abadie, Diamond, and Hainmueller \[2010\]](#)). Per capita cigarette sales (in packs) is the most widely used indicator of smoking prevalence found in the tobacco research literature (Abadie, Diamond, and Hainmueller 2010) and serves here as the aggregate outcome variable under study, measured at the state level from 1970 until 2000 (with 1989 representing the first year of the intervention). The current data file was obtained from the `synth` package ([Abadie, Diamond, and Hainmueller 2014](#)), which originally obtained the cigarette sales data and average retail price of cigarettes from [Orzechowski and Walker \(2005\)](#). Eleven states were discarded from the dataset because of their adoption of some other large-scale tobacco control program at some point during California's intervention period under study between 1989 and 2000, leaving 38 states as potential controls (Abadie, Diamond, and Hainmueller 2010).

4.1 Single-group ITSA

In this example, we use `itsa` to assess the impact of Proposition 99 in reducing California's per capita cigarette sales (in packs), using a single-group design. More specifically, we assess whether the introduction of Proposition 99 resulted in a shift in the level and trend of per capita cigarette sales compared with those of the preintervention period (as described in section 2.1).

First, we load the data and declare the dataset as panel:

```
. use cigsales
. tsset state year
    panel variable:  state (strongly balanced)
    time variable:  year, 1970 to 2000
                  delta: 1 unit
```

Next, we specify a single-group ITSA with California (state number 3 in the study) as the treatment group and 1989 as the start of the intervention, request postintervention trend estimates, and plot the results. The model is estimated using `newey` with one lag:

```
. itsa cigsale, single treat(3) trperiod(1989) lag(1) posttrend figure
    panel variable:  state (strongly balanced)
    time variable:  year, 1970 to 2000
                  delta: 1 unit
```

```
Regression with Newey-West standard errors      Number of obs      =      31
maximum lag: 1                                F( 3,      27) =     331.45
                                              Prob > F          =      0.0000
```

cigsale	Coef.	Newey-West Std. Err.	t	P> t	[95% Conf. Interval]	
_t	-1.779474	.3834188	-4.64	0.000	-2.566184	-.9927632
_x1989	-20.0581	4.724395	-4.25	0.000	-29.75175	-10.36444
_x_t1989	-1.494652	.4368201	-3.42	0.002	-2.390933	-.5983715
_cons	134.0053	4.600271	29.13	0.000	124.5663	143.4442

Postintervention Linear Trend: 1989

Treated: `_b[_t]+_b[_x_t1989]`

Linear Trend	Coeff	Std. Err.	t	P> t	[95% Conf. Interval]	
Treated	-3.2741	0.2688	-12.1803	0.0000	-3.8257	-2.7226

As shown in the regression table, the starting level of the per capita cigarette sales was estimated at 134 packs, and sales appeared to decrease significantly every year prior to 1989 by 1.78 packs ($P < 0.0001$, CI = $[-2.57, -0.99]$). In the first year of the intervention (1989), there appeared to be a significant decrease in per capita cigarette sales of 20.06 packs ($P < 0.0001$, CI = $[-29.75, -10.36]$), followed by a significant decrease in the annual trend of sales (relative to the preintervention trend) of 1.49 packs per capita per year ($P = 0.002$, CI = $[-2.39, -0.60]$). We also see, from the `lincom` estimate produced by specifying `posttrend`, that after the introduction of Proposition 99, per capita

cigarette sales decreased annually at a rate of 3.27 packs (95% CI = $[-3.83, -2.72]$). Figure 2 provides a visual display of these results.




Figure 2. Single-group ITSA with Newey–West standard errors and one lag

To ensure that we fit a model that accounts for the correct autocorrelation structure, we use `actest` (Baum and Schaffer 2013), to test for autocorrelation.

```
. actest, lags(6)
```

Cumby-Huizinga test for autocorrelation

H0: variable is MA process up to order q

HA: serial correlation present at specified lags >q

H0: q=0 (serially uncorrelated) HA: s.c. present at range specified				H0: q=specified lag-1 HA: s.c. present at lag specified			
lags	chi2	df	p-val	lag	chi2	df	p-val
1 - 1	15.242	1	0.0001	1	15.242	1	0.0001
1 - 2	15.255	2	0.0005	2	3.300	1	0.0693
1 - 3	15.325	3	0.0016	3	1.192	1	0.2749
1 - 4	15.896	4	0.0032	4	0.000	1	0.9880
1 - 5	16.057	5	0.0067	5	1.113	1	0.2914
1 - 6	16.078	6	0.0133	6	2.051	1	0.1521

Test allows predetermined regressors/instruments

Test requires conditional homoskedasticity

As shown in the right-side panel of the output table, autocorrelation is present at lag 1 but not at any higher lag orders (up to the six lags tested). Thus our initial model specifying `lag(1)` should correctly account for this autocorrelation.

An alternative approach is to rerun `itsa` specifying the `prais` option, which is inherently designed to fit an AR(1) model. Here we specify `rhotype(tscorr)`, which bases p on the autocorrelation of the residuals, and add robust standard errors.

```
. itsa cigsale, single treat(3) trperiod(1989) replace prais rhotype(tscorr)
> vce(robust)
      panel variable:  state (strongly balanced)
      time variable:  year, 1970 to 2000
      delta: 1 unit

(output omitted)
Prais-Winsten AR(1) regression -- iterated estimates
Linear regression                               Number of obs   =       31
                                                F(3, 27)         =    609.24
                                                Prob > F         =     0.0000
                                                R-squared        =     0.9011
                                                Root MSE        =     2.5964
```

cigsale	Coef.	Semirobust Std. Err.	t	P> t	[95% Conf. Interval]	
_t	-1.843139	.4538631	-4.06	0.000	-2.77439	-.9118892
_x1989	-6.094491	.8840197	-6.89	0.000	-7.90835	-4.280633
_x_t1989	-1.998494	.9191	-2.17	0.039	-3.884332	-.1126568
_cons	128.1931	3.958813	32.38	0.000	120.0703	136.316
rho	.9424635					

```

Durbin-Watson statistic (original)    0.535242
Durbin-Watson statistic (transformed) 1.342728
```

Because the estimates produced using `prais` are transformed, they are not directly comparable with those of `newey`, which are produced using an OLS model. However, these results confirm a significant decrease in the annual trend of sales (relative to the preintervention trend) of 2 packs per capita per year ($P = 0.039$, $CI = [-3.88, -0.11]$). `prais` provides the Durbin–Watson d statistic as an indicator of how well the model corrects for first-order autocorrelation. d can take on values between 0 and 4, and under the null hypothesis, d is equal to 2. Values of d less than 2 suggest positive autocorrelation ($p > 0$), whereas values of d greater than 2 suggest negative autocorrelation ($p < 0$); see [R] **regress postestimation time series**. As discussed previously, there are several more intuitive and flexible tests of autocorrelation; however, none of them can currently be used in conjunction with `prais`.

4.2 Multiple-group ITSA

In this example, we use `itsa` to assess the impact of Proposition 99 in reducing California's per capita cigarette sales (in packs), using a multiple-group design. More specifically, we now compare California's experience with that of the other 38 states in the data file.

```
. itsa cigsale, treat(3) trperiod(1989) lag(1) replace figure
      panel variable:  state (strongly balanced)
      time variable:  year, 1970 to 2000
      delta: 1 unit

Regression with Newey-West standard errors      Number of obs      =      1,209
maximum lag: 1                                F( 7,      1201)    =      364.04
                                              Prob > F           =      0.0000
```

cigsale	Coef.	Newey-West Std. Err.	t	P> t	[95% Conf. Interval]	
_t	-.5477701	.2941289	-1.86	0.063	-1.124834	.0292935
_z	-2.041967	5.75639	-0.35	0.723	-13.33567	9.251731
_z_t	-1.231704	.4641182	-2.65	0.008	-2.142276	-.321131
_x1989	-17.25168	3.815452	-4.52	0.000	-24.73737	-9.765987
_x_t1989	-.5035089	.5252893	-0.96	0.338	-1.534096	.5270778
_z_x1989	-2.806417	5.841839	-0.48	0.631	-14.26776	8.654929
_z_x_t1989	-.9911435	.6657528	-1.49	0.137	-2.297311	.3150244
_cons	136.0472	3.818559	35.63	0.000	128.5554	143.539

As shown in the regression table, the initial mean level difference between California and the remaining states (`_z`) was not significant ($P = 0.723$, $CI = [-13.33, 9.25]$), but the difference in the mean baseline slope (`_z_t`) was significant ($P = 0.008$, $CI = [-2.14, -0.32]$). This is verified upon visual inspection of figure 3: the trajectory of mean cigarette sales for the 38 states appears to rise higher than in California, and that level remains elevated throughout the duration of the observation period. Given this differential pattern of change in the baseline, one could argue that the 38 other states were not comparable with California and, thus, treatment-effect estimates for `_z_x1989` and `_z_x_t1989` may be biased (in the present case, both estimates are not statistically significant). Therefore, this model could be improved by limiting the choice of control groups to only those with similar values on these two variables.




Figure 3. Multiple-group ITSA with Newey–West standard errors and one lag; all 38 “nontreated” states are used for comparison

In the following example, we limit the analysis to only those states that are comparable with California on baseline level and trend of the outcome variable, as described in section 2.2. Comparability in the current context is defined as having a p -value greater than 0.10 on both `_z` and `_z_t`. Three comparison states meet this criteria: Colorado, Idaho, and Montana.

As shown in both the regression table and verified upon visual inspection of figure 4, the treatment group is comparable with controls on both baseline level and trend. While there is no statistically significant treatment effect during the first year of the intervention (`_z_x1989`), there is a statistically significant annual reduction in the pre-post trend compared with that of controls of 1.97 per capita cigarette sales per year ($P = 0.003$, $CI = [-3.26, -0.68]$). Additionally, we see from the `posttrend` output that the treatment group decreased annual cigarette sales in the postintervention period by 3.27 packs, the control group decreased sales over the same period by only 1 pack, and the difference between them is 2.28 packs per capita per year.

```
. itsa cigsale, treat(3) trperiod(1989) contid(4 8 19) lag(1) replace posttrend
> figure
    panel variable:  state (strongly balanced)
    time variable:  year, 1970 to 2000
                delta:  1 unit

Regression with Newey-West standard errors      Number of obs      =      124
maximum lag: 1                                F( 7,      116) =      251.48
                                              Prob > F          =      0.0000
```

cigsale	Coef.	Newey-West Std. Err.	t	P> t	[95% Conf. Interval]	
_t	-1.464503	.3837773	-3.82	0.000	-2.224622	-.7043836
_z	2.046198	6.218666	0.33	0.743	-10.27065	14.36305
_z_t	-.3149707	.5330632	-0.59	0.556	-1.37077	.7408282
_x1989	-13.58866	4.180499	-3.25	0.002	-21.86867	-5.308658
_x_t1989	.4746428	.4992501	0.95	0.344	-.514185	1.463471
_z_x1989	-6.469433	6.185239	-1.05	0.298	-18.72008	5.781212
_z_x_t1989	-1.969295	.6533782	-3.01	0.003	-3.263393	-.6751973
_cons	131.9591	4.355318	30.30	0.000	123.3328	140.5853

Comparison of Linear Postintervention Trends: 1989

```
Treated   : _b[_t] + _b[_z_t] + _b[_x_t1989] + _b[_z_x_t1989]
Controls   : _b[_t] + _b[_x_t1989]
Difference : _b[_z_t] + _b[_z_x_t1989]
```

Linear Trend	Coeff	Std. Err.	t	P> t	[95% Conf. Interval]	
Treated	-3.2741	0.2594	-12.6234	0.0000	-3.7878	-2.7604
Controls	-0.9899	0.2883	-3.4336	0.0008	-1.5608	-0.4189
Difference	-2.2843	0.3878	-5.8905	0.0000	-3.0523	-1.5162

These results highlight the importance of ensuring that treatment and control units are comparable on the preintervention level and trend of the outcome variable when conducting a multiple-group ITSA. As described in section 2.2, an iterative process can be used in which each nontreated group is compared separately with the treatment group. Those groups with p -values greater than a specified threshold on both β_4 and β_5 of 3 can be retained as controls for inclusion in the final model. This approach can be easily extended to other covariates as well; however, if achieving balance on many covariates is an important factor, two alternative approaches to ITSA should be considered: the synthetic controls approach described by [Abadie, Diamond, and Hainmueller \(2010\)](#) and implemented in Stata using the `synth` package ([Abadie, Diamond, and Hainmueller 2014](#)), or the propensity-score weighting technique described by [Linden and Adams \(2011\)](#).




Figure 4. Multiple-group ITSA with Newey–West standard errors and one lag; three states, comparable on the baseline level and trend of the outcome, are used for comparison

4.3 Multiple treatment periods

itsa can accommodate design variations in which the effect of multiple treatment periods are of interest. For example, the researcher may be interested in studying the effects of an intervention that is introduced, withdrawn, and reintroduced, or an intervention that is followed by a separate intervention at a later point in time (see Barlow, Hayes, and Nelson [1984] for many other design alternatives).

For exposition, in the following example we add a fictitious intervention to the cigarette sales data, starting in 1982. We reestimate the single-group ITSA from section 4.1, now with one additional intervention period.


```
. itsa cigsale, single treat(3) trperiod(1982 1989) lag(1) replace posttr
> figure
    panel variable:  state (strongly balanced)
    time variable:  year, 1970 to 2000
                delta:  1 unit

Regression with Newey-West standard errors      Number of obs      =      31
maximum lag: 1                                F( 5,      25) =    657.58
                                              Prob > F      =    0.0000
```

cigsale	Coef.	Newey-West Std. Err.	t	P> t	[95% Conf. Interval]	
_t	-.2038463	.3343824	-0.61	0.548	-.8925198	.4848273
_x1982	-8.040472	2.875704	-2.80	0.010	-13.96309	-2.117849
_x_t1982	-3.639012	.378855	-9.61	0.000	-4.419278	-2.858745
_x1989	-9.285162	2.561198	-3.63	0.001	-14.56005	-4.010275
_x_t1989	.5687319	.3153918	1.80	0.083	-.0808296	1.218293
_cons	125.2333	2.281397	54.89	0.000	120.5347	129.932

Postintervention Linear Trend: 1982

Treated: _b[_t]+_b[_x_t1982]

Linear Trend	Coeff	Std. Err.	t	P> t	[95% Conf. Interval]	
Treated	-3.8429	0.2237	-17.1757	0.0000	-4.3037	-3.3821

Postintervention Linear Trend: 1989

Treated: _b[_t]+_b[_x_t1982]+_b[_x_t1989]

Linear Trend	Coeff	Std. Err.	t	P> t	[95% Conf. Interval]	
Treated	-3.2741	0.2793	-11.7205	0.0000	-3.8495	-2.6988

The interpretation of all coefficients up to the second intervention is as before. That is, the first intervention period is compared with the preintervention period. However, the additional coefficients for the second intervention period, `_x1989` and `_x_t1989`, are now compared with those of the prior (first) intervention period.

As shown in both the regression table and verified upon visual inspection of figure 5, there is evidence of a “treatment effect” beginning in 1982 and no additional decrease in annual sales after the implementation of the second intervention (which in reality was the true intervention period).




Figure 5. Single-group ITSA with Newey–West standard errors and two intervention periods

We can demonstrate this effect further via the `posttrend` option, which estimates the postintervention trends separately after the first and second intervention periods. As shown in the `posttrend` output, the annual decrease in cigarette sales after 1982 was 3.84 packs per year, while the annual decrease in sales after 1989 was slightly less, at 3.27 packs per year (the difference is 0.569, which appears in the original regression table as `_x_t1989`). Thus the results of this exercise reveal an additional utility of defining multiple treatment periods in `itsa` when analyzing single-group data: it allows for the testing of “interruptions” away from the true start of the intervention, as described in section 2.1. As demonstrated in this example, a statistically significant reduction in the trend of cigarette sales in California began several years prior to implementation of Proposition 99. This result also further highlights the importance, even when using ITSA, of finding a comparable control group to represent the counterfactual.

5 Discussion

While the randomized controlled trial remains the gold standard research design, there are situations in which this design is not feasible or practical, such as when large-scale interventions or policy changes target the entire population. When data are available for multiple time points in both the preintervention and the postintervention periods, interrupted time-series designs offer a robust quasi-experimental alternative for evaluating treatment effects (Campbell and Stanley 1966; Shadish, Cook, and Campbell 2002).

In this article, I have demonstrated the basic implementation of `itsa` to estimate treatment effects for a single treatment group, a multiple-group comparison, and when more than one intervention has been employed sequentially. Additional important issues were also addressed, such as criteria for choosing and specifying a model, testing for autocorrelation, robustness testing for interruptions prior to the true intervention start period, and choosing comparable controls. More-complex models can easily be estimated with `itsa` by including additional covariates to control for confounding, seasonal effects, and the impact of external events. Moreover, the addition of key time-series variables to the dataset after running `itsa` allows for further estimation of treatment effects using more complex OLS models or `arima`, assuming that more sophisticated time-series modeling is warranted and assuming the availability of a sufficient number of observations.

6 Acknowledgments

I thank Nicholas J. Cox for his support while developing `itsa`, and I thank Steven J. Samuels for creating the `posttrend` option and for assistance with several post-review changes to `itsa`. I also thank Michael J. Harvey for his assistance with \LaTeX , and Steven J. Samuels, Roger B. Newson, John L. Adams, Andrew Ryan, and Julia Adler-Milstein for their reviews and helpful comments on the article. I also thank the anonymous reviewer and chief editor for their thoughtful reviews and recommendations for improving both the article and the `itsa` command.

7 References

- Abadie, A., A. Diamond, and J. Hainmueller. 2010. Synthetic control methods for comparative case studies: Estimating the effect of California's tobacco control program. *Journal of the American Statistical Association* 105: 493–505.
- . 2014. `synth`: Stata module to implement synthetic control methods for comparative case studies. Statistical Software Components S457334, Department of Economics, Boston College. <https://ideas.repec.org/c/boc/bocode/s457334.html>.
- Barlow, D. H., S. C. Hayes, and R. O. Nelson. 1984. *The Scientist Practitioner: Research and Accountability in Clinical and Educational Settings*. New York: Pergamon Press.
- Baum, C. F., and M. E. Schaffer. 2013. `actest`: Stata module to perform Cumby–Huizinga general test for autocorrelation in time series. Statistical Software Components S457668, Department of Economics, Boston College. <https://ideas.repec.org/c/boc/bocode/s457668.html>.
- Biglan, A., D. Ary, and A. C. Wagenaar. 2000. The value of interrupted time-series experiments for community intervention research. *Prevention Science* 1: 31–49.
- Box, G. E. P., and G. M. Jenkins. 1976. *Time Series Analysis: Forecasting and Control*. San Francisco, CA: Holden Day.

- Box, G. E. P., and G. C. Tiao. 1975. Intervention analysis with applications to economic and environmental problems. *Journal of the American Statistical Association* 70: 70–79.
- Briesacher, B. A., S. B. Soumerai, F. Zhang, S. Toh, S. E. Andrade, J. L. Wagner, A. Shoaibi, and J. H. Gurwitz. 2013. A critical review of methods to evaluate the impact of FDA regulatory actions. *Pharmacoepidemiology and Drug Safety* 22: 986–994.
- Campbell, D. T., and J. C. Stanley. 1966. *Experimental and Quasi-Experimental Designs for Research*. Chicago, IL: Rand McNally.
- Cochrane Effective Practice and Organisation of Care (EPOC) Group. 2013. EPOC-specific resources for review authors. Oslo: Norwegian Knowledge Centre for the Health Services.
<http://epocoslo.cochrane.org/epoc-specific-resources-review-authors>.
- Crosbie, J. 1993. Interrupted time-series analysis with brief single-subject data. *Journal of Consulting and Clinical Psychology* 61: 966–974.
- Cumby, R. E., and J. Huizinga. 1992. Testing the autocorrelation structure of disturbances in ordinary least squares and instrumental variables regressions. *Econometrica* 60: 185–195.
- Gillings, D., D. Makuc, and E. Siegel. 1981. Analysis of interrupted time series mortality trends: An example to evaluate regionalized perinatal care. *American Journal of Public Health* 71: 38–46.
- Glass, G. V., V. L. Willson, and J. M. Gottman. 1975. *Design and Analysis of Time-Series Experiments*. Boulder, CO: Colorado Associated University Press.
- Gottman, J. M. 1981. *Time-series Analysis: A Comprehensive Introduction for Social Scientists*. New York: Cambridge University Press.
- Huitema, B. E., and J. W. McKean. 2000a. Design specification issues in time-series intervention models. *Educational and Psychological Measurement* 60: 38–58.
- . 2000b. A simple and powerful test for autocorrelated errors in OLS intervention models. *Psychological Reports* 87: 3–20.
- Imbens, G. W., and T. Lemieux. 2008. Regression discontinuity designs: A guide to practice. *Journal of Econometrics* 142: 615–635.
- Kutner, M. H., C. J. Nachtsheim, J. Neter, and W. Li. 2005. *Applied Linear Statistical Models*. New York: McGraw-Hill.
- Linden, A., and J. L. Adams. 2011. Applying a propensity-score based weighting model to interrupted time series data: Improving causal inference in program evaluation. *Journal of Evaluation in Clinical Practice* 17: 1231–1238.

- McDowall, D., R. McCleary, E. E. Meidinger, and R. A. Hay. 1980. *Interrupted Time Series Analysis*. Newbury Park, CA: Sage Publications.
- McKnight, S., J. W. McKean, and B. E. Huitema. 2000. A double bootstrap method to analyze linear models with autoregressive error terms. *Psychological Methods* 5: 87–101.
- Muller, A. 2004. Florida's motorcycle helmet law repeal and fatality rates. *American Journal of Public Health* 94: 556–558.
- Orzechowski, W., and R. C. Walker. 2005. *The Tax Burden on Tobacco. Historical Compilation*, vol. 40. Arlington, VA: Orzechowski & Walker.
- Ramsay, C. R., L. Matowe, R. Grilli, J. M. Grimshaw, and R. E. Thomas. 2003. Interrupted time series designs in health technology assessment: Lessons from two systematic reviews of behavior change strategies. *International Journal of Technology Assessment in Health Care* 19: 613–623.
- Riley, W. T., R. E. Glasgow, L. Etheredge, and A. P. Abernethy. 2013. Rapid, responsive, relevant (R3) research: A call for a rapid learning health research enterprise. *Clinical and Translational Medicine* 2: 1–6.
- Shadish, W. R., T. D. Cook, and D. T. Campbell. 2002. *Experimental and Quasi-Experimental Designs for Generalized Causal Inference*. Boston: Houghton Mifflin.
- Simonton, D. K. 1977a. Cross-sectional time-series experiments: Some suggested statistical analyses. *Psychological Bulletin* 84: 489–502.
- . 1977b. Erratum to Simonton. *Psychological Bulletin* 84: 1097.
- Velicer, W. F., and J. Harrop. 1983. The reliability and accuracy of time series model identification. *Evaluation Review* 7: 551–560.
- Velicer, W. F., and R. P. McDonald. 1991. Cross-sectional time series designs: A general transformation approach. *Multivariate Behavioral Research* 26: 247–254.

About the author

Ariel Linden is a health services researcher specializing in the evaluation of health care interventions. He is both an independent consultant and an adjunct associate professor at the University of Michigan in the department of Health Management and Policy, where he teaches program evaluation.